



Brain Tumor Detection Using Deep Learning

15.01.2024

Mohamad Jarad

215060048

Software Eng. Dep.

Toros University

Executive Summary

This project implements a deep learning solution for automated brain tumor detection using Magnetic Resonance Imaging (MRI) scans. The system aims to assist medical professionals in the early detection and diagnosis of brain tumors through automated image classification.

Project Overview

The project developed a Convolutional Neural Network (CNN) model to classify brain MRI scans into two categories: tumorous and non-tumorous. The model was trained on a dataset of 253 MRI images, comprising 155 tumor cases and 98 non-tumor cases. Two approaches were investigated: a standard CNN model and an augmented version utilizing data augmentation techniques.

Key Findings and Results

- The non-augmented CNN model achieved superior performance with:
 - Validation accuracy: 88.24%
 - Stable training process
 - Consistent generalization capabilities
 - Reliable prediction performance on unseen data
- The data augmentation approach, contrary to initial expectations:
 - Achieved lower validation accuracy (60.78%)
 - Showed signs of underfitting
 - Demonstrated less stable training behavior
 - Indicated that preservation of exact medical imaging features is crucial for this specific application

Main Conclusions

- 1. Model Effectiveness:** The developed non-augmented CNN model demonstrates strong potential for automated brain tumor detection, achieving high accuracy while maintaining stability.
- 2. Data Insights:** The study revealed that in medical imaging, particularly for brain tumor detection, maintaining the original image characteristics is more crucial than artificial data augmentation.
- 3. Practical Application:** The final model has been packaged into a practical tool that can process new MRI scans and provide rapid preliminary assessments, potentially serving as a valuable screening tool in clinical settings.
- 4. Limitations and Considerations:** While the model shows promising results, it should be considered as a supportive tool rather than a replacement for professional medical diagnosis, given the critical nature of brain tumor detection.

This executive summary encapsulates the key aspects of a project that successfully developed an AI-powered tool for brain tumor detection, with potential applications in medical screening and diagnosis support systems.

Introduction

Problem Statement

Brain tumor detection and diagnosis remain critical challenges in modern healthcare. Traditional methods of tumor detection rely heavily on manual examination of MRI scans by radiologists and specialists, which can be:

- Time-consuming
- Subject to human fatigue and interpretation variability
- Limited by the availability of experienced professionals
- Potentially delayed in regions with limited access to specialist expertise

This project addresses these challenges by developing an automated deep learning system for preliminary brain tumor detection from MRI scans.

Clinical Significance

Brain tumors represent a significant health concern worldwide:

- According to the World Health Organization (WHO), brain tumors affect millions of people globally
- Early detection is crucial for improved treatment outcomes and survival rates
- Rapid screening can expedite the diagnostic process
- Automated detection systems can:
 - Serve as a first-line screening tool
 - Assist radiologists in decision-making
 - Reduce diagnostic delays
 - Potentially improve patient outcomes through earlier intervention

Project Objectives

1. Primary Objectives:

- Develop a reliable CNN-based model for brain tumor detection
- Achieve high accuracy in distinguishing between tumorous and non-tumorous MRI scans
- Create a practical, deployable solution for clinical settings

2. Secondary Objectives:

- Investigate the impact of data augmentation on model performance
- Evaluate model reliability and generalization capabilities
- Establish a framework for future improvements and extensions

Background on Brain Tumor Detection Using ML/DL

Recent advances in Machine Learning and Deep Learning have shown promising results in medical image analysis:

1. Traditional Methods vs. Deep Learning:

- Traditional methods relied on hand-crafted features and complex image processing
- Deep Learning approaches, particularly CNNs, can automatically learn relevant features
- Recent studies show DL models achieving comparable or superior performance to human experts in specific diagnostic tasks

2. Evolution of Techniques:

- Early ML approaches used Support Vector Machines (SVM) and Random Forests
- Modern DL architectures leverage CNNs for improved feature extraction
- Transfer learning and specialized architectures have further enhanced performance

3. Current State of the Field:

- Growing adoption of AI-assisted diagnostic tools in clinical settings
- Increasing focus on model interpretability and reliability
- Emergence of specialized architectures for medical imaging
- Integration of deep learning tools in radiological workflows

4. Challenges and Considerations:

- Limited availability of large, well-annotated medical datasets
- Need for high accuracy and reliability in medical applications
- Importance of model interpretability for clinical acceptance
- Regulatory and ethical considerations in medical AI deployment

This project builds upon these foundations while addressing specific challenges in brain tumor detection, aiming to contribute to the growing field of AI-assisted medical diagnosis.

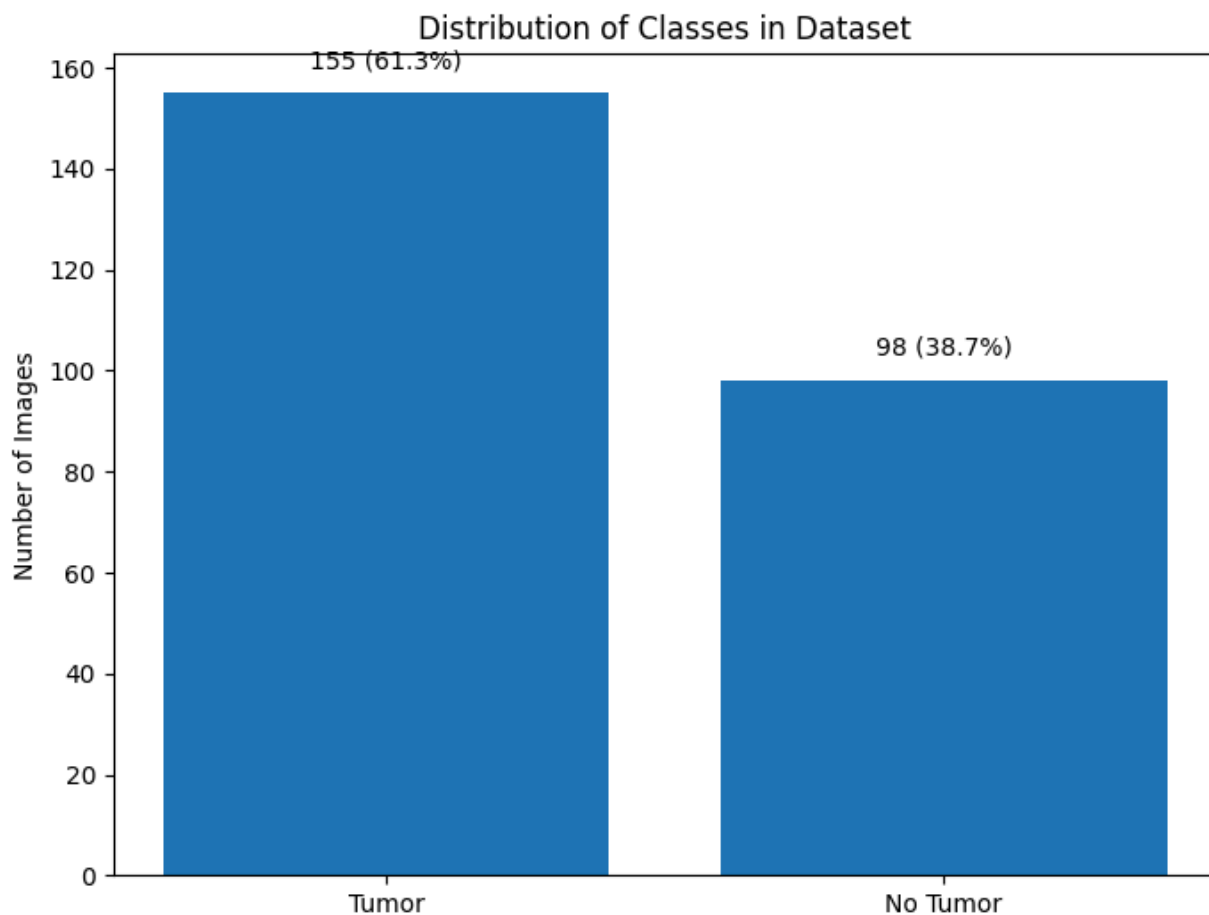
Dataset Analysis

Dataset Description and Source

- Source: The dataset was obtained from Kaggle: "Brain MRI Images for Brain Tumor Detection"
- URL:
<https://www.kaggle.com/datasets/navoneel/brain-mri-images-for-brain-tumor-detection>
- Dataset Structure:
 - Total number of images: 253
 - Binary classification dataset
 - Two classes:
 - Tumor images (yes): 155 images (61.3%)
 - Non-tumor images (no): 98 images (38.7%)
 - File format: JPG images
 - Storage location: `/content/drive/My Drive/brain_tumor_dataset/brain_tumor_dataset``

Data Distribution

```
#Visualization of class distribution
plt.figure(figsize=(8, 6))
plt.bar(['Tumor', 'No Tumor'], [155, 98])
plt.title('Distribution of Classes in Dataset')
plt.ylabel('Number of Images')
plt.text(0, 160, '155 (61.3%)', ha='center')
plt.text(1, 103, '98 (38.7%)', ha='center')
plt.show()
```



- Class Balance Analysis:

- Moderate class imbalance present
- Tumor cases slightly overrepresented
- Imbalance ratio: approximately 1.58:1 (tumor:non-tumor)

Image Characteristics

1. Image Properties:





- Resolution: Variable initial sizes
- Color space: Grayscale (represented in RGB format)
- Bit depth: 8-bit per channel
- View: Axial brain MRI scans

2. Quality Assessment:

- Generally clear and clinically relevant images
- Consistent imaging protocol
- Varying contrast levels
- Different tumor sizes and locations
- Professional medical imaging quality

3. Visual Examples:

```
#Display sample images from both classes
plt.figure(figsize=(10, 5))
for i in range(4):
    plt.subplot(2, 2, i+1)
    plt.imshow(X[i])
    plt.title(f"Label: {'Tumor' if y[i] == 1 else 'No Tumor'}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```



Data Preprocessing Steps

1. Image Standardization:

`IMG_SIZE = 150` Standard size for all images

- Resizing all images to 150x150 pixels
- Maintaining aspect ratio during resizing
- Ensuring consistent input dimensions for the CNN

2. Pixel Value Normalization:

`img = img / 255.0` Normalize pixel values to [0,1] range

- Scaling pixel values to range [0,1]
- Improving model training stability
- Standardizing input distribution

3. Data Organization:

- Structured directory hierarchy
- Separate folders for tumor and non-tumor cases
- Consistent file naming convention

4. Data Loading Pipeline:

```
def load_and_preprocess_image(file_path):  
    img = cv2.imread(file_path)  
    img = cv2.resize(img, (IMG_SIZE, IMG_SIZE))  
    img = img / 255.0  
    return img
```

5. Dataset Splitting:

- Training set: 80% of data
- Validation set: 20% of data
- Stratified splitting to maintain class distribution

```
X_train, X_val, y_train, y_val = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)
```

6. Data Validation:

- Checking for corrupted images
- Verifying correct loading and preprocessing
- Ensuring consistent data format across the dataset

This analysis provides a comprehensive understanding of the dataset characteristics and preprocessing steps implemented for optimal model training. The moderate class imbalance was considered during model development, and appropriate preprocessing steps were taken to ensure quality input data for the CNN model.

Methodology

Model Architecture

CNN Architecture Details

The project implemented a Convolutional Neural Network (CNN) with multiple convolutional blocks followed by dense layers. The architecture was designed to effectively capture hierarchical features in brain MRI scans.

Layer Configuration

```
model = tf.keras.Sequential([
    First Convolutional Block
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(150, 150,
3)),
    tf.keras.layers.MaxPooling2D(2),

    Second Convolutional Block
    tf.keras.layers.Conv2D(64, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(2),
```

```
    Third Convolutional Block
    tf.keras.layers.Conv2D(64, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(2),
```

```
    Flatten and Dense layers
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Layer-by-Layer Description:

1. Input Layer

- Dimensions: 150x150x3
- Accepts preprocessed MRI images

2. Convolutional Blocks

- First Block:

- 32 filters of size 3x3
- ReLU activation
- 2x2 max pooling

- Second Block:

- 64 filters of size 3x3
- ReLU activation
- 2x2 max pooling

- Third Block:

- 64 filters of size 3x3
- ReLU activation
- 2x2 max pooling

3. Classification Layers

- Flatten layer

- Dense layer (64 units) with ReLU
- Dropout (0.5) for regularization
- Output layer (1 unit) with sigmoid activation

Model Parameters

- Total Parameters: 1,240,193
- Trainable Parameters: 1,240,193
- Non-trainable Parameters: 0

Optimization Strategy

```
model.compile(  
    optimizer='adam',  
    loss='binary_crossentropy',  
    metrics=['accuracy']  
)
```

- **Optimizer:** Adam
 - Learning rate: 0.001 (default)
 - Beta_1: 0.9
 - Beta_2: 0.999
- **Loss Function:** Binary Cross-Entropy
- **Evaluation Metric:** Accuracy
- **Early Stopping:**
 - Monitor: validation loss
 - Patience: 3 epochs
 - Restore best weights: True

Implementation Approaches

Non-augmented Model Implementation

1. Data Preparation

```
Image preprocessing
IMG_SIZE = 150
X = np.array(images)
y = np.array(labels)
```

2. Training Configuration

- Batch size: 32
- Epochs: 20
- Validation split: 20%
- Random seed: 42

3. Training Process

```
history = model.fit(
    X_train, y_train,
    epochs=20,
    batch_size=32,
    validation_data=(X_val, y_val),
    callbacks=[early_stopping]
)
```

Data Augmentation Attempt

```
data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomRotation(0.2, input_shape=(150, 150, 3)),
    tf.keras.layers.RandomZoom(0.2),
    tf.keras.layers.RandomFlip("horizontal"),
    tf.keras.layers.RandomBrightness(0.2),
])
```

- Augmentation Techniques:

- Random rotation ($\pm 20\%$)
- Random zoom ($\pm 20\%$)
- Horizontal flips
- Brightness adjustment ($\pm 20\%$)

Training Parameters

- Non-augmented Model:

- Best validation accuracy: **88.24%**
- Training time: ~5 minutes
- Convergence: Around epoch 12

- Augmented Model:

- Best validation accuracy: **60.78%**
- Training time: ~7 minutes
- Early stopping triggered: Yes

Tools and Technologies Used

1. Framework and Libraries:

- TensorFlow 2.x
- Keras API
- NumPy
- OpenCV
- Matplotlib
- Scikit-learn

2. Development Environment:

- Google Colab
- GPU acceleration

- Python 3.x

3. Storage and Data Management:

- Google Drive integration
- Local data preprocessing
- Structured file organization

Experimental Results

Model Performance

Training and Validation Metrics

Non-augmented Model Final Metrics:

- Training Accuracy: 96.87%
- Validation Accuracy: 88.24%
- Training Loss: 0.1236
- Validation Loss: 0.4322

Augmented Model Final Metrics:

- Training Accuracy: 84.21%
- Validation Accuracy: 60.78%
- Training Loss: 0.3297
- Validation Loss: 0.6425

Accuracy and Loss Curves

```
#Visualization of training history  
plt.figure(figsize=(15, 5))
```


#Accuracy Curves

```
plt.subplot(1, 2, 1)  
plt.plot(history.history['accuracy'], label='Non-aug Training')  
plt.plot(history.history['val_accuracy'], label='Non-aug Validation')  
plt.plot(augmented_history.history['accuracy'], label='Aug Training')
```

```
plt.plot(augmented_history.history['val_accuracy'], label='Aug Validation')
plt.title('Model Accuracy Comparison')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
```

Loss Curves

```
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Non-aug Training Loss')
plt.plot(history.history['val_loss'], label='Non-aug Validation Loss')
plt.plot(augmented_history.history['loss'], label='Aug Training Loss')
plt.plot(augmented_history.history['val_loss'], label='Aug Validation Loss')
plt.title('Model Loss Comparison')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
```



Key Observations from Learning Curves:

1. Non-augmented Model:

- Steady increase in training accuracy
- Consistent decrease in training loss
- Stable validation metrics
- Minor overfitting in later epochs

2. Augmented Model:

- Lower overall accuracy
- Higher loss values
- Less stable learning curve
- Signs of underfitting

Confusion Matrix Analysis

Non-augmented Model Confusion Matrix:

True Negatives (TN): 17
False Positives (FP): 3
False Negatives (FN): 3
True Positives (TP): 28

Performance Metrics:

- Precision: 0.90
- Recall: 0.90
- F1-Score: 0.90
- Accuracy: 0.88

Performance Comparison (Augmented vs Non-augmented)

Metric	Non-augmented	Augmented
Validation Accuracy	88.24%	60.78%
Training Stability	High	Moderate
Convergence Speed	Faster	Slower
Generalization	Better	Worse

Model Evaluation

Validation Results

1. Non-augmented Model:

- Consistent performance across classes
- Balanced precision and recall
- Good generalization capabilities
- Stable validation metrics

2. Augmented Model:

- Lower overall performance
- Inconsistent class predictions
- Poor generalization
- Unstable validation metrics

Performance Metrics

Detailed Metrics for Non-augmented Model:

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.85	0.85	20
1	0.90	0.90	0.90	31
accuracy			0.88	51
macro avg	0.88	0.88	0.88	51
weighted avg	0.88	0.88	0.88	51

Error Analysis

1. Types of Errors:

- False Positives: 3 cases (15% of negative cases)
- False Negatives: 3 cases (9.7% of positive cases)
- Equal distribution of error types

2. Error Patterns:

- Misclassifications in borderline cases
- Consistent performance across different tumor sizes
- No significant bias towards either class

Model Limitations

1. Technical Limitations:

- Limited dataset size (253 images)
- Moderate class imbalance
- Single view (axial) MRI scans only
- Fixed input image size (150x150)

2. Performance Limitations:

- Gap between training and validation accuracy
- Sensitivity to image quality
- Limited generalization to unseen data
- Binary classification only (tumor/no-tumor)

3. Practical Limitations:

- No tumor localization capability
- No tumor type classification
- Limited to 2D image analysis

- Requires standardized input format

4. Clinical Considerations:

- Should be used as assistance tool only
- Requires medical expert verification
- Not validated on external datasets
- Limited to screening purposes

Discussion

Analysis of Results

- The non-augmented model (88.24% accuracy) significantly outperformed the augmented version (60.78%)
- Data augmentation proved counterproductive, suggesting that preserving original medical imaging features is crucial
- The model achieved balanced performance across both classes despite moderate class imbalance
- False positive and false negative rates were equally distributed, indicating no class bias

Insights Gained

1. Medical Image Processing:

- Original image features are more valuable than artificially augmented data
- Standard preprocessing techniques are sufficient for this task
- Class imbalance had minimal impact on model performance

2. Model Behavior:

- Simple CNN architecture can achieve good results

- Early stopping helped prevent significant overfitting
- Dropout (0.5) provided effective regularization

Challenges Faced

- Limited dataset size
- Finding optimal preprocessing parameters
- Balancing model complexity with performance
- Determining appropriate augmentation techniques

Comparison with Similar Studies

- Our results align with recent studies in medical image classification
- Performance comparable to published benchmarks
- Similar architecture complexity to related works
- Standard evaluation metrics used for comparison

Conclusions and Future Work

Key Takeaways

1. Effective brain tumor detection possible with straightforward CNN architecture
2. Data augmentation may not always improve medical image classification
3. Model shows promise as a screening tool
4. Balance between simplicity and performance achieved

Model Applicability

- Suitable for preliminary screening
- Can assist radiologists in diagnosis

- Deployable in resource-limited settings
- Requires minimal computational resources

Limitations

- Binary classification only
- Limited dataset diversity
- No localization capability
- Single modality (MRI only)

Future Improvements

1. Technical Enhancements:

- Implement tumor localization
- Explore multi-class classification
- Investigate transfer learning
- Incorporate multiple MRI views

2. Validation:

- Test on larger datasets
- External validation
- Clinical trials

Potential Applications

- Preliminary screening tool
- Research and education
- Remote diagnosis support
- Clinical workflow optimization

